biggy: An Implementation of Unified Architecture for Big Data Management System

Yao Wu

Email: ideamaxwu@ruc.edu.cn

Abstract—**Big data is emerging and reaching every corner of industries and every filed of researches. Big data can be interpreted as a term for large or complex datasets that traditional data processing applications are inadequate. Challenges in big data include data analysis, data transaction, data query, data storage, data visualization, data privacy and so on. Various tools and systems are proposed and developed to face these challenges on different emphases. In this paper, we propose datar, a new prospective and unified architecture for Big Data Management System (BDMS) from the point of system architecture by leveraging ideas from computers. We introduce five key components of datar by reviewing the current status of BDMS. Moreover, the architecture of datar is presented as an implementation, i.e., biggy. Manipulation details of biggy are demonstrated by several examples. Our work shows the envisioned datar as a feasible solution and unified architecture that can manage big data pluggably, automatically and intelligently with specific functionalities, where specific functionalities refer to input, storage, computation, control and output of big data in this paper.**

I. INTRODUCTION

A. From Computer to Datar

As Turing proposed the question "Can machines think?" [1], the imitation game begins. However, before that, Von Neumann had started an engineering research on computer and described an logical design of a computer using the stored-program concept in 1945, which has controversially come to be known as the Von Neumann architecture [2]. More ago, English mathematician and computer pioneer Charles Babbage proposed the Analytical Engine, a mechanical general-purpose computer designed. The goal of these pioneers is to design a computing machine better than human brain, which can liberate the tedious computation from manual work.

During the last several decades, data management principles such as physical and logical independence, declarative querying and cost-based optimization have led to several fields of researches and a prosperous industry. More importantly, these technical advances have enabled the first round of business intelligence applications and laid the foundation for managing and analyzing big data today. Many novel challenges and opportunities associated with big data necessitate rethinking many aspects of these data management platforms, while retaining other desirable aspects. The practice [3] and theory [4] contribution of Bachman and Codd open up the research on database. And the steps on the road to data management never stop such as, Ingres [5], Postgres [6], Mariposa [7], C-Store [8] and VoltDB [9].

Since the computing power of machines becomes stronger, we can sniff the shift from computation to data management to explore more in-sight information and knowledge from
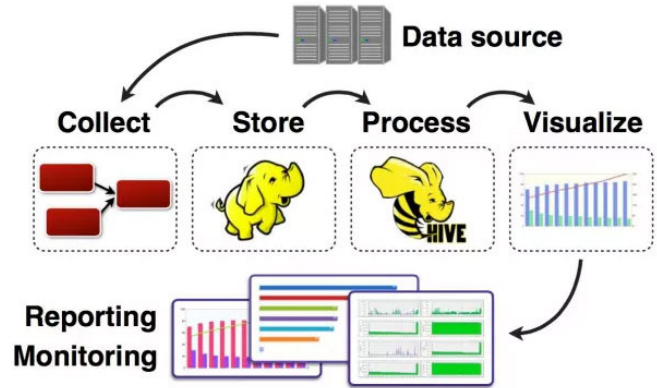

Fig. 1. A typical workflow for big data management.

data. From computers to data management systems, we never stop liberating human from tedious and complex work. As we all know, (Big Data Management System) BDMS is a complex set of functionalities, we think it necessary to propose an unified architecture to guide the design of BDMS. From these observations, we summarize and describe several main components in BDMS to provide a full understanding for better explanation of our proposed datar architecture.

As shown in Fig. 1, BDMS consists of several core components such as to, collect, storage, process and visualize. Compared with traditional database systems, BDMS architecture is more flexible and open due to various focus-ons. In this paper, we unify the BDMS as **datar**, a proposed general architecture to design and build BDMS, with respect to the term **computer**.

As we all knnow, the popular computer architecture is divided into five parts, input, storage, computation, control and output, in which, computation is the center. If we look closely, we can find that, the BDMS is just the same as computer, consisting of (data) input, (data) storage, (data) computation (query/analysis), (data) control (transaction) and data output (visualization), in which, data storage is the center. In other words, we can call a computer Fast Computation Processing System (FCPS). Likewise, the BDMS is called as datar, a BDMS centered at data. To better explain the similarities and differences between a computer and a datar, in terms of architecture, we use Fig. 2 to illustrate. In Fig. 2 (a), five core components of a computer are shown in separate rectangles, while in Fig. 2 (b), five corresponding parts are shown. A computer and a datar share the similar functionalities with different emphases on computation and data storage.
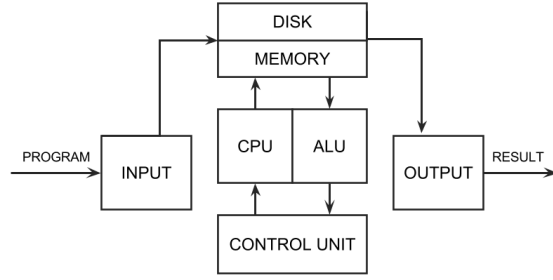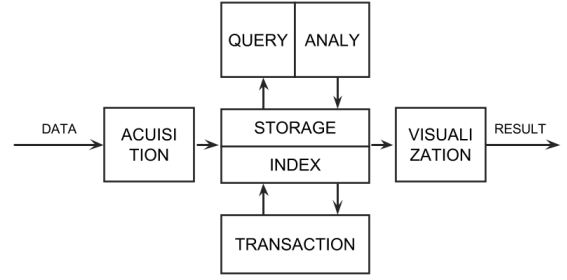
(a) Computer Architecture
(b) Datar Architecture

Fig. 2. Computer VS Datar comparison of architecture.

B. What is Datar

Definition (Datar) *A datar is a stack of coherent softwares with specific functionalities that can be applied to mine valuable information from persisted data automatically, where specific functionalities refer to input, storage, computation, control and output of big data.* In this paper, we implement datar as biggy based on AsterixDB [10] with its extensions, Spark-MLlib and d3 as an materialization of this envisioned architecture. biggy is a data storage centered solution for datar implementation.

A datar, i.e., a full-function BDMS, consists of five parts, data input, data storage, data computation, data control and data output. Compared to the computation-centered computer, a datar is data-centered. We take AsterixDB for example, which is a new, full-function BDMS. Data input is how the data gets into the system. In AsterixDB, data feed is a built-in mechanism allowing new data to be continuously ingested into system from external sources, incrementally populating the datasets and their associated indexes [11]. Data storage is how the data is stored in the system and how the indexes are built. In AstrixDB, data and index are stored based on LSM structure [12]. Data computation is how to mine valuable information from stored data. A bunch of methods can be applied, such as popular in-memory computation framework, Spark. Besides, the execution of data processing is also part of data analysis, like Hyracks [13] in AsterixDB. Data control is how to control data when it is processed. It is different from the traditional database systems which have strict ACID. Another very important aspect of datar is data output, e.g., visualization. Cloudberry[1] is a research prototype to support interactive analytics and visualization of large amounts of spatial-temporal data using AsterixDB. Based on these features of AsterixDB, it is ideal for us to build biggy on top of it and extend it with more functions like complex analysis by Spark and data visualization by d3.

---

[1]http://cloudberry.ics.uci.edu/

C. How Far Can Datar Go

In the late 1970s, the concept of "database machine" emerged, which is a technology specially used for storing and analyzing data. With the increase of data volume, the storage and processing capacity of a single mainframe computer system became inadequate. In the 1980s, people proposed "share nothing", a parallel database system, to meet the demand of the increasing data. The share nothing system architecture is based on the use of cluster and every machine has its own processor, storage, and disk. In the late 1990s, the advantages of parallel database was widely recognized in the database field.

However, many challenges facing big data arise. With the development of Internet services, indexes and queried contents were rapidly growing. Therefore, search engine companies had to face the challenges of handling such big data. Google created storage systems GFS and programing models MapReduce to cope with the challenges brought about by data management and analysis at the Internet scale. In January 2007, Jim Gray, a pioneer of database software, called such transformation "The Fourth Paradigm". He also thought the only way to cope with such paradigm was to develop a new generation of computing tools to manage, visualize, and analyze massive data.

Over the past few years, nearly all major companies, including EMC, Oracle, IBM, Microsoft, Google, Amazon, and Facebook, etc. have started their big data projects. Big data shows great value in real application and challenges arise. At present, data has become an important production factor that could be comparable to material assets and human capital. However, some literature discuss obstacles in the development of big data applications, such as data representation, data reduction and compression and data life cycle management.

In this paper, we describe the Big Data Management System (BDMS) from a new perspective, the view of a computer, to propose datar. We focus our attention on the system architecture in the BDMS and break it down into several parts to elaborate. The envisioned datar is implemented as biggy. The key contributions can be summarized as,

- We review current big data management systems as datar by five core components and analyze the advantages and disadvantages.

- An unified architecture for big data management, i.e., datar, is proposed and explained to manage big data pluggablly, automatically and intelligently.

- We implement the envisioned architecture as biggy based on AsterixDB with several other systems to fulfill the functions of input, storage, computation, control and output, and demonstrate biggy in details.

The paper is organized as follows. In Section II, we introduce the BDMS as datar from five aspects, input, storage, control, computation and output. Section III introduces the framework and implementation of dater, biggy. The details of manipulation of biggy are given in Section IV. Finally, we conclude with futures in section V.

II. BREAK DOWN BDMS

A. Data Input

Data can be input into storage since its generation or be "inserted" from other resources.

1) Data Generation: **Web data** is unstructured data from the web, also including social network data, location based service data and the linked pages. Web Data is very complex as compared to traditional text documents. **Enterprise Data**, the internal data of enterprises are the main sources of big data. The internal data of enterprises mainly consists of online trading data and online analysis data, most of which are historically static data and are managed by traditional relation DBMSs in a structured manner. **Government Data** is collected from government agencies. Inundated with database schemes, documents, emails, web content and XML, agencies are left with daunting integration challenges. **Other Data** generates from more sources. As scientific applications are increasing, the scale of datasets is gradually expanding, and the development of some disciplines greatly relies on the analysis of masses of data. In addition, pervasive sensing and computing among natural, commercial and social environments are generating heterogeneous data with unprecedented complexity. These datasets have their unique data characteristics in scale, time dimension, and data category.

2) Data Feed: Data feed is also known as data acquisition, having continuous data arrive into DBMS from external sources and incrementally populate a persisted dataset and associated indexes. In the past, ETL (Extract Transform Load) systems work the same way. It is also similar to streaming data processing. A simple way of having data being put into a big data management system on a continuous basis is to have a single program fetch data from an external data source, parse the data, and then invoke an insert statement per record or batch of records. It is hard to reason about the data consistency, scalability and fault-tolerance offered by an assembly 'gluing' together different systems. Therefore, it is natural for a BDMS to provide "native" support for data feed management.

Flume[2] is a distributed, reliable, and available service for efficiently collecting, aggregating, and moving large amounts

Key-Value	Column	Document	Graph	Other (NewSQL)
Dynamo	BigTable	MongoDB	Giraph	SAP HANA
Voldemort	Cassandra	SimpleDB	Neo4j	TiDB
Redis	Hbase	CouchDB	OrientDB	VoltDB
MemcacheDB	HyperTable		Pregel	NuoDB
Scalaris	C-store		FlockDB	Google Spanner
TiKV				

TABLE I.   POPULAR BIG DATA STORAGE SYSTEM.

of log data. It fetches data from sources through channels to sinks. Kafka[3] is used for building real-time data pipelines and streaming applications in categories called topics from producers to consumers. In AsterixDB, we build a fault-tolerant data feed facility that scales through partitioned parallelism by using a high-level language. A generic plug-and-play model can help datar cater to a wide variety of data sources and applications.

B. Data Storage

Data storage focuses on raw data storage and indexes storage in this paper. Besides, schemes storage, configuration files storage and views storage are also data storage but out of our scope.

1) Data Storage: Big data storage refers to the storage and management of large-scale datasets while achieving reliability and availability of data accessing. Various storage systems emerge to meet the demands of massive data. There are four main NoSQL databases technologies: key-value DB, column-oriented DB, document-oriented DB and graph-oriented DB.

Key-Value DB Key-value databases are constituted by a simple data model and data is stored corresponding to key-values. Every key is unique and customers may input queried values according to the keys. Such databases feature a simple structure and the modern key-value databases are characterized with high expandability and shorter query response time than those of relational databases. Examples are Dynamo[4] [14] by Amazon,Voldemort[5] [15] by LinkedIn and Redis[6] [16], Memcache DB[7] [17], Scalaris[8] [18], Riak and Tokyo Tyrant.

Column-oriented DB The column-oriented databases store and process data according to columns other than rows. Both columns and rows are segmented in multiple nodes to realize expandability. BigTable[9] [19] by Google, Cassandra[10] [20] by Facebook, Hbase[11] [21] cloned from BigTable, HyperTable[12] are main column-based databases.

Documnet-oriented DB Compared with key-value storage, document storage can support more complex data forms. Since documents do not follow strict modes, there is no need to

---

[2]https://flume.apache.org/

[3]https://kafka.apache.org/

[4]https://aws.amazon.com/dynamodb/

[5]http://www.project-voldemort.com/

[6]http://redis.io/

[7]http://memcachedb.org/

[8]http://scalaris.zib.de/

[9]https://cloud.google.com/bigtable/

[10]http://cassandra.apache.org/

[11]http://hbase.apache.org/
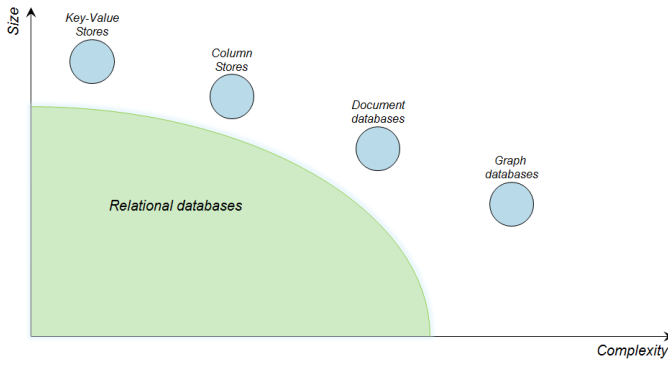
[12]http://www.hypertable.org/

Fig. 3. Comparision of data storage model.

conduct mode migration. Examples are MongoDB[13] [22] open-source, SimpleDB[14] by Amazon, CouchDB[15] [23] on Apache.

Graph-oreinted DB A graph database is a database that uses graph structures for semantic queries with nodes, edges and properties to represent and store data. Apache Giraph[16], Neo4J[17], OrientDB[18] are designed and implemented for purpose of graph storage.

In addition, other companies and researchers also have their solutions to meet the different demands for storage of big data. Table I shows summary of the popular big data storage systems and Fig. 3 shows the size and complexity comparison among these different types of data management systems. One example is Pluggable Storage Engines (e.g., MySQL) provided by Voldemort. TiDB also supports external database connection. Further more, **NewSQL** is a class of modern relational database management systems that seek to provide the same scalable performance of NoSQL systems for online transaction processing (OLTP) read-write workloads while still maintaining the ACID guarantees of a traditional database system. In this situation, storage is just like disk that can be added, replaced and deleted.

2) Data Index: Index is always an effective method to reduce the expense of disk reading and writing, and improve insertion, deletion, modification, and query speeds in both traditional relational databases that manage structured data, and other technologies that manage semistructured and un-structured data. However, index has a disadvantage that it has the additional cost for storing index files which should be maintained dynamically when data is updated.

Basic structures include Hash table, Tree-based index, Multidimensional index, and Bitmap index. Big data index [24] has additional requirements, such as parallelism and easily partitioned into pieces for parallel processing. Artificial Intelligence Indexing approaches are so called because of their ability to detect unknown behavior in big data. They establish relationships between data items by observing patterns and categorizing items or objects with similar traits. Latent Semantic Indexing and Hidden Markov Model are two popular

AI indexing approaches. In Non-AI indexing approaches, the formation of indexes does not depend on the meaning of the data item or the relationship between texts. Rather, indexes are formed based on items most queried or searched for in a particular data set.

C. Data Computation

Data computation can be from simple SQL-like query to complex machine learning techniques. There is no clear boundary to divide them, but we introduce them by two coarse categories as data query and data analysis.

1) Data Query: MapReduce [25], Dryad[19] [26], All-Pairs, Pregel [27], Spark[20] [28] are the popular programming models and execution engines. More in-memory databases are proposed to accelerate the computation.

MapReduce is a programming model and an associated implementation for processing and generating large data sets with a parallel, distributed algorithm on a cluster. The Dryad Project is investigating programming models for writing parallel and distributed programs to scale from a small cluster to a large data-center. Pregel is Google's scalable and fault-tolerant platform with an API that is sufficiently flexible to express arbitrary graph algorithms. Apache Spark is a fast and general engine for big data processing, with built-in modules for streaming, SQL, machine learning and graph processing.

2) Data Analysis: The analysis can be from simple statistic to deep data mining technology. Nowadays, deep learnig has become a trend in analysis of big data. MLlib[21] [29] by Java, Scipy, Theano, Caffe[22] [30] by Python, TensorFlow[23] [31] by C++, Torch, etc.

MLlib is Spark's machine learning library, focusing on learning algorithms and utilities, including classification, regression, clustering, collaborative filtering, dimensionality reduction, as well as underlying optimization primitives. Caffe is a deep learning framework made with expression, speed, and modularity in mind. TensorFlow is an open source software library for machine learning in various kinds of perceptual and language understanding tasks.

D. Data Control

According to CAP theorem, it is not feasile for BDMS to fullfil ACID (Atomic Consistent Isolation Durable) from traditional databases. However, BASE (Basic Availability Soft-state Eventual consistency) is an alternative. CAP theorem says it is impossible for a protocol to guarantee both consistency and availability in a partition prone distributed system as shown in Fig. 4. Most NoSQL database system architectures favor one factor over the other.

BigTable, used by Google App engine, and HBase, which runs over Hadoop, claim to be strongly consistent within a data-center and highly available meaning there is an eventual consistency between data-centers. Updates are propagated to

[13] https://www.mongodb.com/
[14] https://aws.amazon.com/simpledb/
[15] http://couchdb.apache.org/
[16] http://giraph.apache.org/
[17] https://neo4j.com/
[18] http://orientdb.com/

[19] https://www.microsoft.com/research/project/dryad/
[20] http://spark.apache.org/
[21] http://spark.apache.org/mllib/
[22] http://caffe.berkeleyvision.org/
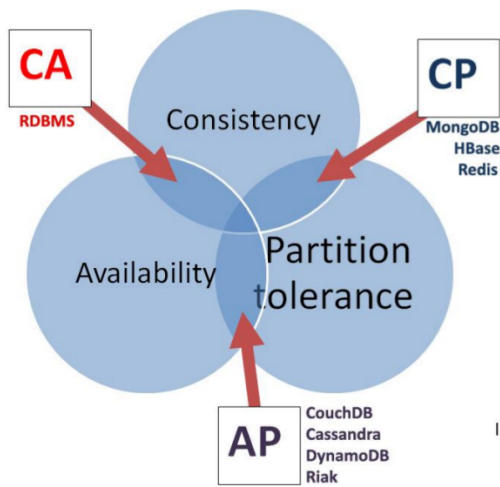[23] https://www.tensorflow.org/

Fig. 4. CAP theorem with associated NoSQL DBs.

all replicas asynchronously. Amazon's Dynamo, Cassandra and Riak instead sacrifice consistency in favor of availability and partition tolerance. They achieve a weaker form of consistency known as eventual consistency, without guarantees on the order of updates across replicas and when they will be applied.

1) Data Transaction: In databases, a transaction is a set of separate actions that must all be completely processed, or none processed at all. In partitioned databases, trading some consistency for availability can lead to dramatic improvements in scalability. It is hard to leverage between the speed and scale of NoSQL, and the relational, transactional strength and consistency of traditional RDBMS.

NewSQL [32] is a kind of next-generation scalable relational database management systems for Online Transaction Processing (OLTP) that provide scalable performance of NoSQL systems for read-write workloads, as well as maintaining the ACID guarantees of a traditional database system. One of the famous NewSQL is Google Spanner [33], which is a globally distributed NewSQL database.

NuoDB[24] is a distributed database designed with SQL service: all the properties of ACID transactions, standard SQL language support and relational logic. ClustrixDB[25] is a distributed SQL database built for large-scale and fast-growing applications. VoltDB [9] is an insanely fast in-memory database with incredible high read and write speeds. CouchDB read operations use a Multi-Version Concurrency Control (MVCC) model.

Although NewSQL systems vary greatly in their internal architectures, the two distinguishing features common among them is that they all support the relational data model and use SQL as their primary interface

2) Data Recovery: Big data applications as well as operational systems must be supported by a robust and rapid recovery process. As database architecture has fundamentally changed to meet new application requirements, data protection needs to be redefined and re-architected as well.

Big data shook up the database arena, ushering in a new class of "scale out" technologies. The scale-out nature of the architecture can also be difficult for traditional backup applications to handle. Organizations that are deploying big data platforms and applications must realize the importance of backing up their data. Platform-provided mechanisms such as replicas and snapshots are not sufficient to ensure proper data protection and to minimize downtime. Proper backup and recovery requires some investment but is well worth it given the role big data plays in driving business value. Some of the most common mechanisms[26] include:

- Multiple replicas of data eliminates the need for separate backup/recovery tools of big data.

- Lost data can be quickly and easily rebuilt from the original raw data.

- Backing up a petabyte of big data is not economical or practical.

- Remote disaster recovery copies can serve as a backup copy.

- Writing backup/recovery scripts for big data is easy.

- Big data backup/recovery operations costs are very small.

- Snapshots are an effective backup mechanism for big data.

3) Resource Management: Big data computation always runs on thousands of machines, which needs the resource management among clusters. Mesos [34] is a platform for sharing commodity clusters between multiple diverse cluster computing frameworks. The fundamental idea of Hadoop YARN [35] is to split up the functionalities of resource management and job scheduling/monitoring into separate daemons. Apache ZooKeeper [36] is an effort to develop and maintain an open-source server which enables highly reliable distributed coordination.

E. Data Output

Visualization is the best way to present the results of big data management. Besides, in this section, we also mention data sharing as a way of data output.

1) Data Visualization: Visualization helps us take a deep look into the big data, which provides us a interactive and graphic way to embrase the inside of big data. Tools like Tableau[27], Plotly[28], Visual.ly[29] are emerging. Zeppelin[30] is a web-based notebook that enables interactive data analytics.

Big data analytics plays a key role through reducing the data size and complexity in big data applications. Visualization is an important approach to helping big data get a complete view of data and discover data values. Big data analytics

---

[24]http://www.nuodb.com/

[25]http://www.clustrix.com/

[26]http://www.networkworld.com/article/3113036/big-data-business-intelligence/debunking-the-most-common-big-data-backup-and-recovery-myths.html

[27]http://www.tableau.com/

[28]https://plot.ly/

[29]http://visual.ly/

[30]https://zeppelin.apache.org/

and visualization should be integrated seamlessly so that they work best in big data applications. Many conventional data visualization methods are often used, such as table, histogram, scatter plot, time line, data flow diagram, and entity relationship diagram, etc. Visualizations are not only static; they can be interactive. Interactive visualization can be performed through approaches such as zooming. Scalability and dynamics are two major challenges in visual analytics [37]. Visualization of big data with diversity and heterogeneity (structured, semi-structured, and unstructured) is a big problem. Big data visualization can be performed through a number of approaches such as more than one view per representation display, dynamical changes in number of factors and filtering.

2) Data Sharing: Because sharing data [38] will necessarily increase the potential benefit to society of the subject's participation by providing greater opportunities for scientific discovery, Brakewood and Poldrack argued that researchers may have an ethical duty to share their data unless doing so would increase risk to the subjects.

In research filed, many insititutions share their data to promote the research. Governments also start share data to public for common benefits. However, some data like personal data and enterprise internal data cannot share since its privacy and confidetial. Therefore, some guidelines our regulations should be made to lead us properly share data, rather than share all or share nothing.

III. BUILD UP DATAR

A. Datar Hypothesis

As we have discussed, we envision an universal architecture of big data management systems, **datar**, to make one more step. The idea essentially comes from computer. A datar is a set of coherent softwares/systems that can manage big data pluggablly, automatically and intelligently with specific functionalities. Pluggibility means any of the five parts can be replaced by one another easily, automation means the flow from data input to data output can be executed coherently and intelligence means the ultimate goal of datar is to mine valuable information in depth.

B. biggy Framework and Implementation

To put the envisioned datar into practice, we implement it as **biggy**. biggy is based on AsterixDB with several other systems to fulfill the functions of input, storage, computation, control and output. Currently, we plan to implement biggy based on AsterixDB, BAD, Spark-MLlib and d3. AsterixDB is the core component for data storage and control, BAD for data input, Spark-MLlib for data computation and d3 for data output. We make it more pluggable and automatic rather than just gluing them together.The project can be found at Github[31].

Fig. 5 shows the implementation framework of biggy. As we can see, AsterixDB provides LSM-based data storage with B^+-tree and R-tree indexes structure, 2PL concurrency control and write-ahead-log (WAL) recovery strategy. For data computation, AsterixDB SQL++ works the data query and Spark MLlib for data analysis. Data input is based on AsterixDB
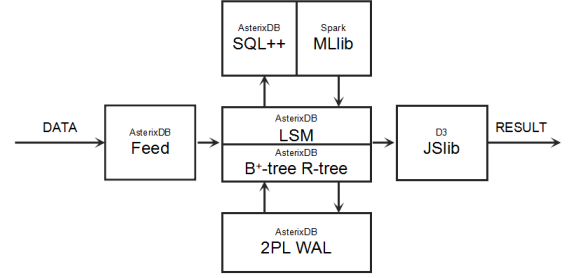

Fig. 5. Framework of biggy.

Fig. 6. Relations among several key terms

feed feature and data output based on d3 JSlib. Further work of supporting most popular systems (e.g., TensorFlow) as plugins needs done with fulfillment of intelligence.

IV. PLAY WITH BIGGY

A. biggy Installation and Management

For better explanation, relations among several key terms are shown in Fig. 6. bigo is an instance of biggy and bigdb is the physical files where store the data. Fig. 7 shows the configuration of biggy. Before installation and use of biggy, it is necessary to set the five parts in configuration file. For each module, there should be a adapter to make it fit in biggy, which is implemented by programmers and can be shared for public to use. That is why we make it open source and any modules/systems/libraries follow the design of biggy can be plugged in to work. Fig. 8 shows how to install and manage biggy, in which bigo is the instance of biggy. Management includes *create, start, use, stop, delete and describe*.

B. biggy Manipulation of Instance bigo

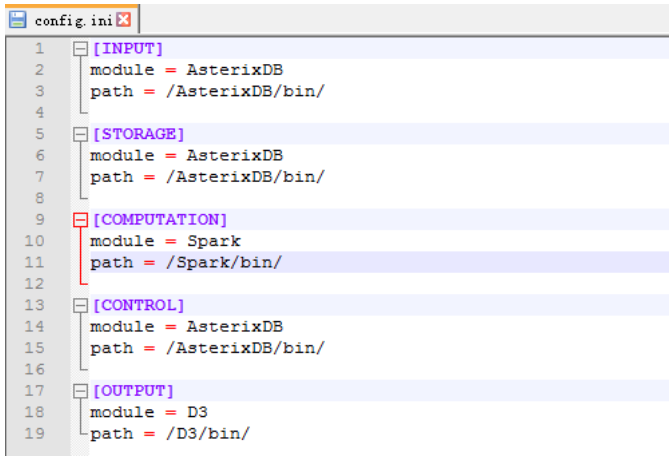Fig. 9 shows the manipulation of bigo, input, storage, computation, control and output.

---

[31]https://github.com/Ideamaxwu/biggy

```ini
config.ini ✕
1   [INPUT]
2   module = AsterixDB
3   path = /AsterixDB/bin/
4
5   [STORAGE]
6   module = AsterixDB
7   path = /AsterixDB/bin/
8
9   [COMPUTATION]
10  module = Spark
11  path = /Spark/bin/
12
13  [CONTROL]
14  module = AsterixDB
15  path = /AsterixDB/bin/
16
17  [OUTPUT]
18  module = D3
19  path = /D3/bin/
```

Fig. 7. Configuration of biggy.

- install biggy project

```
./bigo.sh install biggy
```

- manage [create, start, use, stop, delete, describe] biggy instance named bigo

```
./bigo.sh [new, start, use, stop, delete, info] biggy bigo
```
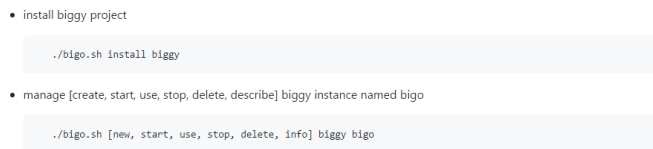
Fig. 8. Example of biggy installation and management.

1) bigo Input: Data INPUT is to have continous data arrive into biggy from external sources and incrementally populate a persisted dataset and associated indexes. The external data source can be files, data in another database or streaming data fed by a feed.

2) bigo Storage: Data STORAGE is to store the data on physical storage devices. A new database can be created in this procedure.

3) bigo Computation: Data COMPUTATION is to explore valuable information from persisted data by simple query, data mining and deep learning. Basic data query like insert, delete and update. Complex analysis are provided by Spark MLlib.

4) bigo Control: Data CONTROL is to make all the data management conducted without conflicts, faults or failures. Currently, it is a builtin module in AsterixDB, which can not be modified. Other implementation of datar can be applied with other data control options. For biggy, it is relatively dependent on the core base of AsterixDB.

5) bigo Output: Data OUTPUT provides a way to show the results. In biggy, we apply d3 as the visualization tool. An example is shown in Fig. 10.
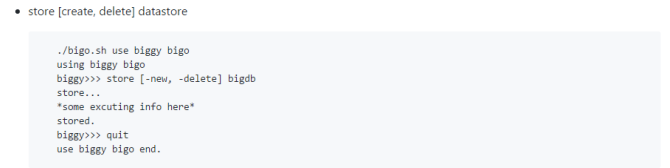
V. CONCLUSION AND FUTURE WORK

We have illustrated the datar as an envisioned architecture for BDMS from the perspective of a computer in five components and propose biggy as an implementation of datar to manage big data pluggablly, automatically and intelligently with specific functionalities.
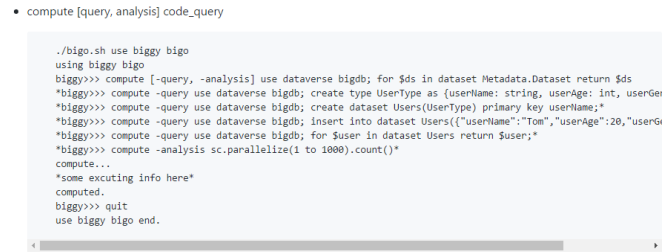
We have entered an era of Big Data. Though better analysis of the large volumes of data that are becoming available, there is the potential for making faster advances

- input [feed, file] path_from path_to

```
*TBD module*
./bigo.sh use biggy bigo
using biggy bigo
biggy>>> input [-feed, -file] data/source.xls data/destination.table
input...
*some excuting info here*
inputed.
biggy>>> quit
use biggy bigo end.
```

(a) Input

- store [create, delete] datastore

```
./bigo.sh use biggy bigo
using biggy bigo
biggy>>> store [-new, -delete] bigdb
store...
*some excuting info here*
stored.
biggy>>> quit
use biggy bigo end.
```

(b) Storage

- compute [query, analysis] code_query

```
./bigo.sh use biggy bigo
using biggy bigo
biggy>>> compute [-query, -analysis] use dataverse bigdb; for $ds in dataset Metadata.Dataset return $ds
*biggy>>> compute -query use dataverse bigdb; create type UserType as {userName: string, userAge: int, userGe
*biggy>>> compute -query use dataverse bigdb; create dataset Users(UserType) primary key userName;*
*biggy>>> compute -query use dataverse bigdb; insert into dataset Users({"userName":"Tom","userAge":20,"userGe
*biggy>>> compute -query use dataverse bigdb; for $user in dataset Users return $user;*
*biggy>>> compute -analysis sc.parallelize(1 to 1000).count()*
compute...
*some excuting info here*
computed.
biggy>>> quit
use biggy bigo end.
```

(c) Computation

- control

```
*BUILDIN module*
./bigo.sh use biggy bigo
using biggy bigo
biggy>>> control
control...
Built-in Module. NO Configuration!
controlled.
biggy>>> quit
use biggy bigo end.
```

(d) Control

- output [visual, share] path_data

```
./bigo.sh use biggy bigo
using biggy bigo
biggy>>> output -visual
output...
*some excuting info here*
outputed.
biggy>>> quit
use biggy bigo end.
```

(e) Output

Fig. 9. Example of bigo manipulation.

in many scientific disciplines and improving the profitability and success of many enterprises. However, many technical challenges must be addressed before this potential can be realized fully. The challenges include not just the obvious issues of scale, but also heterogeneity, lack of structure, error-handling, privacy, timeliness, provenance, and visualization, at all stages of the analysis pipeline from data acquisition to result interpretation. Furthermore, these challenges will require transformative solutions, and will not be addressed naturally by the next generation of industrial products. We must support and encourage fundamental research towards addressing these architecture design and technical challenges if we are to achieve the promised benefits of Big Data.
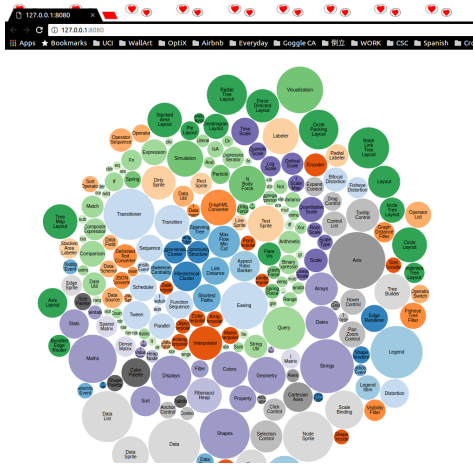
Fig. 10.   Example of bigo visualization by d3.

REFERENCES

[1] A. M. Turing, "Computing machinery and intelligence," *Mind*, vol. 59, no. 236, pp. 433–460, 1950.

[2] J. Von Neumann, "First draft of a report on the edvac," *IEEE Annals of the History of Computing*, no. 4, pp. 27–75, 1993.

[3] C. W. Bachman, "On a generalized language for file organization and manipulation," *Commun. ACM*, vol. 9, no. 3, pp. 225–226, 1966.

[4] E. F. Codd, "A relational model of data for large shared data banks," *Commun. ACM*, vol. 13, no. 6, pp. 377–387, 1970.

[5] M. Stonebraker, E. Wong, P. Kreps, and G. Held, "The design and implementation of INGRES," *ACM Trans. Database Syst.*, vol. 1, no. 3, pp. 189–222, 1976.

[6] M. Stonebraker, "The postgres DBMS," in *Proceedings of the 1990 ACM SIGMOD International Conference on Management of Data, Atlantic City, NJ, May 23-25, 1990.*, 1990, p. 394.

[7] M. Stonebraker, P. M. Aoki, W. Litwin, A. Pfeffer, A. Sah, J. Sidell, C. Staelin, and A. Yu, "Mariposa: A wide-area distributed database system," *VLDB J.*, vol. 5, no. 1, pp. 48–63, 1996.

[8] M. Stonebraker, D. J. Abadi, A. Batkin, X. Chen, M. Cherniack, M. Ferreira, E. Lau, A. Lin, S. Madden, E. J. O'Neil, P. E. O'Neil, A. Rasin, N. Tran, and S. B. Zdonik, "C-store: A column-oriented DBMS," in *Proceedings of the 31st International Conference on Very Large Data Bases, Trondheim, Norway, August 30 - September 2, 2005*, 2005, pp. 553–564.

[9] M. Stonebraker and A. Weisberg, "The voltdb main memory DBMS," *IEEE Data Eng. Bull.*, vol. 36, no. 2, pp. 21–27, 2013.

[10] S. Alsubaiee, Y. Altowim, H. Altwaijry, A. Behm, V. Borkar, Y. Bu, M. Carey, I. Cetindil, M. Cheelangi, K. Faraaz *et al.*, "Asterixdb: A scalable, open source bdms," *Proceedings of the VLDB Endowment*, vol. 7, no. 14, pp. 1905–1916, 2014.

[11] R. Grover and M. J. Carey, "Data ingestion in asterixdb." in *EDBT*, 2015, pp. 605–616.

[12] S. Alsubaiee, A. Behm, V. Borkar, Z. Heilbron, Y.-S. Kim, M. J. Carey, M. Dreseler, and C. Li, "Storage management in asterixdb," *Proceedings of the VLDB Endowment*, vol. 7, no. 10, pp. 841–852, 2014.

[13] V. Borkar, M. Carey, R. Grover, N. Onose, and R. Vernica, "Hyracks: A flexible and extensible foundation for data-intensive computing," in *2011 IEEE 27th International Conference on Data Engineering*. IEEE, 2011, pp. 1151–1162.

[14] G. DeCandia, D. Hastorun, M. Jampani, G. Kakulapati, A. Lakshman, A. Pilchin, S. Sivasubramanian, P. Vosshall, and W. Vogels, "Dynamo: amazon's highly available key-value store," *ACM SIGOPS Operating Systems Review*, vol. 41, no. 6, pp. 205–220, 2007.

[15] R. Sumbaly, J. Kreps, L. Gao, A. Feinberg, C. Soman, and S. Shah, "Serving large-scale batch computed data with project voldemort," in *Proceedings of the 10th USENIX conference on File and Storage Technologies*. USENIX Association, 2012, pp. 18–18.

[16] J. L. Carlson, *Redis in Action*. Manning Publications Co., 2013.

[17] B. Fitzpatrick, "Distributed caching with memcached," *Linux journal*, vol. 2004, no. 124, p. 5, 2004.

[18] T. Schütt, F. Schintke, and A. Reinefeld, "Scalaris: reliable transactional p2p key/value store," in *Proceedings of the 7th ACM SIGPLAN workshop on ERLANG*. ACM, 2008, pp. 41–48.

[19] F. Chang, J. Dean, S. Ghemawat, W. C. Hsieh, D. A. Wallach, M. Burrows, T. Chandra, A. Fikes, and R. E. Gruber, "Bigtable: A distributed storage system for structured data," *ACM Transactions on Computer Systems (TOCS)*, vol. 26, no. 2, p. 4, 2008.

[20] A. Lakshman and P. Malik, "Cassandra: a decentralized structured storage system," *ACM SIGOPS Operating Systems Review*, vol. 44, no. 2, pp. 35–40, 2010.

[21] L. George, *HBase: the definitive guide*. " O'Reilly Media, Inc.", 2011.

[22] K. Banker, *MongoDB in action*. Manning Publications Co., 2011.

[23] J. C. Anderson, J. Lehnhardt, and N. Slater, *CouchDB: the definitive guide*. " O'Reilly Media, Inc.", 2010.

[24] F. B. Adamu, A. Habbal, S. Hassan, R. Les Cottrell, B. White, and I. Abdullahi, "A survey on big data indexing strategies."

[25] J. Dean and S. Ghemawat, "Mapreduce: simplified data processing on large clusters," *Communications of the ACM*, vol. 51, no. 1, pp. 107–113, 2008.

[26] M. Isard, M. Budiu, Y. Yu, A. Birrell, and D. Fetterly, "Dryad: distributed data-parallel programs from sequential building blocks," in *ACM SIGOPS Operating Systems Review*, vol. 41, no. 3. ACM, 2007, pp. 59–72.

[27] G. Malewicz, M. H. Austern, A. J. Bik, J. C. Dehnert, I. Horn, N. Leiser, and G. Czajkowski, "Pregel: a system for large-scale graph processing," in *Proceedings of the 2010 ACM SIGMOD International Conference on Management of data*. ACM, 2010, pp. 135–146.

[28] M. Zaharia, M. Chowdhury, M. J. Franklin, S. Shenker, and I. Stoica, "Spark: cluster computing with working sets." *HotCloud*, vol. 10, pp. 10–10, 2010.

[29] X. Meng, J. Bradley, B. Yuvaz, E. Sparks, S. Venkataraman, D. Liu, J. Freeman, D. Tsai, M. Amde, S. Owen *et al.*, "Mllib: Machine learning in apache spark," *JMLR*, vol. 17, no. 34, pp. 1–7, 2016.

[30] Y. Jia, E. Shelhamer, J. Donahue, S. Karayev, J. Long, R. Girshick, S. Guadarrama, and T. Darrell, "Caffe: Convolutional architecture for fast feature embedding," in *Proceedings of the 22nd ACM international conference on Multimedia*. ACM, 2014, pp. 675–678.

[31] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin *et al.*, "Tensorflow: Large-scale machine learning on heterogeneous distributed systems," *arXiv preprint arXiv:1603.04467*, 2016.

[32] M. Stonebraker, "Newsql: An alternative to nosql and old sql for new oltp apps," *Communications of the ACM. Retrieved*, pp. 07–06, 2012.

[33] J. C. Corbett, J. Dean, M. Epstein, A. Fikes, C. Frost, J. J. Furman, S. Ghemawat, A. Gubarev, C. Heiser, P. Hochschild *et al.*, "Spanner: Googles globally distributed database," *ACM Transactions on Computer Systems (TOCS)*, vol. 31, no. 3, p. 8, 2013.

[34] B. Hindman, A. Konwinski, M. Zaharia, A. Ghodsi, A. D. Joseph, R. H. Katz, S. Shenker, and I. Stoica, "Mesos: A platform for fine-grained resource sharing in the data center." in *NSDI*, vol. 11, no. 2011, 2011, pp. 22–22.

[35] V. K. Vavilapalli, A. C. Murthy, C. Douglas, S. Agarwal, M. Konar, R. Evans, T. Graves, J. Lowe, H. Shah, S. Seth *et al.*, "Apache hadoop yarn: Yet another resource negotiator," in *Proceedings of the 4th annual Symposium on Cloud Computing*. ACM, 2013, p. 5.

[36] P. Hunt, M. Konar, F. P. Junqueira, and B. Reed, "Zookeeper: Wait-free coordination for internet-scale systems." in *USENIX annual technical conference*, vol. 8, 2010, p. 9.

[37] L. Wang, G. Wang, and C. A. Alexander, "Big data and visualization: methods, challenges and technology progress," *Digital Technologies*, vol. 1, no. 1, pp. 33–38, 2015.

[38] R. A. Poldrack and K. J. Gorgolewski, "Making big data open: data sharing in neuroimaging," *Nature neuroscience*, vol. 17, no. 11, pp. 1510–1517, 2014.